

BOB REVIEW ONLY. Do not send this PDF to Hogan. He asked for bullets, then for a DOI under each point. Table 2 and his HKO paragraph are unchanged. The live manuscript is not cut. Confirmatory freeze remains a human decision.

Mechanism points for Bob's review

Nine September 2026. Seven points from the sendable list, each at review length. Hypotheses, refusals, and neighbours are labelled as such. Count ratios below are copied from the live identification article. No new Hospital Authority model is fitted here.

The estimand of the parent paper is an ecological monthly count ratio. Separate negative-binomial models of first hospitalisation after a first diagnosis of coronary heart disease (CHD; 156,156 events) or heart failure (HF; 29,681 events), among people with type 2 diabetes and/or hypertension in Hong Kong, January 2013–December 2023 (132 months). Calendar-month indicators, a 4-df time-trend spline, and a days-in-month offset. Admission cause is absent. Stroke is named and not delivered. Official hot nights, very hot days, and cold days are reported per five additional days in the month. That five is a linear reporting scale, not a consecutive-day trigger.

Point 1. Interrupted overnight recovery after hot nights

This is the keep-mechanism for the CHD residual. It is a hypothesis about why a monthly official hot-night *count* can differ from monthly mean temperature. It is not a sleep coefficient, not a mediator test, and not a finding of Model 1.

Under Newey–West lag-6 reporting, the CHD count ratio is 1.022 (1.002–1.042) per five official hot nights, with Benjamini–Hochberg $q = 0.192$. The three continuous temperature encodings for CHD sit near the null: 0.993 (0.976–1.010) for monthly mean temperature, 0.993 (0.979–1.007) for mean maximum, and 0.994 (0.977–1.012) for mean minimum. CHD very hot days are 0.999 (0.974–1.025). The residual that is being interpreted is therefore specific to the official night *count*, not to “heat” as a single monthly average. All twelve Model 1 q -values exceed 0.19. The CHD hot-night interval additionally depends on the standard-error construction: model-based and HC1 intervals include 1. The pre-2020 window is 1.011 (0.991–1.032) and includes 1. Those facts stay in the Results. Overnight recovery does not repair them.

The scientific content of the hypothesis is narrow. Sleep onset in humans is coupled to a nighttime fall in core temperature. If that fall is interrupted, sleep is shorter and more fragmented, and overnight cardiovascular recovery is a plausible casualty. A month that contains many official hot nights ($T_{\min} \geq 28^\circ\text{C}$ at the Hong Kong Observatory Headquarters) can therefore differ from a month that is merely warm on average, even if the two months have similar mean temperature. Equation (1) sees only the monthly official count X_t . It does not see sleep, core temperature, or indoor temperature.


Chevance, Minor, Vielma and colleagues (2024) collected that sleep evidence in a systematic review of ambient heat and sleep. Higher outdoor or indoor temperature is associated with shorter sleep and poorer sleep quality; fragmentation is among the quality indicators in that literature. The review is about sleep. It is not first CHD hospitalisation in a diabetes or hypertension cohort, not Hong Kong, and not a monthly official-count model. It earns its place as the reason a night *count* is a different question from a monthly mean. It does not identify the 1.022.

Ashe, Wozniak, Conner and colleagues (2025) are the limit on that hypothesis, not a second keep-mechanism. Their scoping review of extreme heat, sleep, and cardiovascular measures found three eligible studies. Blood pressure did not move in one direction. No study demonstrated that sleep mediates a heat–hospitalisation path. The live Discussion already uses this paper as a refusal: the panel cannot treat sleep as a measured mediator. Keep Chevance for the sleep association. Keep Ashe so the overnight-recovery sentence cannot be misread as a dipping theorem.

The local epidemiological neighbour is Guo, Chan, Qiu, Wong and Ho (2024), not Chevance. In daily unplanned emergency hospitalisations in Hong Kong during the hot seasons of 2000–2019, the official hot-night *flag* (minimum temperature $\geq 28^\circ\text{C}$) showed no overall association with non-cancer, non-external hospitalisation over lags 0–4 days after adjustment for multi-day mean temperature (excess relative risk -0.2% , -1.2% to 0.7%). An hourly nighttime excess-heat metric for 20:00–07:59 was associated, including 3.1% ($1.5\text{--}4.8\%$) at its extreme value and 3.4% ($0.2\text{--}6.8\%$) for circulatory admissions. Those percentages are Guo’s daily estimand. They are not this paper’s monthly count ratio, not first CHD after a first diagnosis, and not 2013–2023 as the sole window. They do show, in this city, that an official night encoding and a nighttime intensity encoding can disagree. Monthly official *counts* are coarser than either daily encoding. The contrast between 1.022 and 0.993 is therefore a contrast between monthly questions. It is not a replication of Guo’s hourly result, and it is not evidence that a hot-night effect was “missed” by the means.

What this panel cannot do follows directly. It cannot see sleep. It cannot see indoor temperature. It cannot see overnight recovery. It only counts nights. The first-event totals also fall by about half across the window while official hot nights rise from 10 in 2013 to 61 in 2021. The same 4-df spline absorbs both. That is why the CHD residual is trend-sensitive at the stiff end (3-df 1.011, 0.990–1.032, includes 1) and why overnight recovery must remain a labelled hypothesis rather than a physiological confirmation of Table 2.


Keep the mechanism in the draft, in the live wording: interrupted overnight recovery is one hypothesis for why a hot-night count could differ from monthly mean temperature; the present monthly design cannot test that hypothesis. Do not add a sleep pathway diagram to the manuscript. Do not import Guo’s 3.1% as a magnitude. Do not freeze a confirmatory night claim from this review.



REAL public HKO Headquarters. Official hot night: $T_{\min} \geq 28^\circ\text{C}$. Official cold day: $T_{\min} \leq 12^\circ\text{C}$.

Why a hot-night count can differ from monthly mean temperature

Equation (1) sees X_t ,
not sleep



SCHEMATIC. Not a fitted model. Official hot night = $T_{min} \geq 28^\circ\text{C}$ at HKO Headquarters.

Point 2. Cold-related afterload on a failing heart

This is the keep-mechanism for the HF residual. It is haemodynamic rather than nocturnal. It is a hypothesis about why official cold days can be a different monthly question from monthly mean temperature for a first heart-failure hospitalisation. It is not a blood-pressure coefficient, not an afterload measurement, and not an imported daily relative risk.

Under Newey–West lag-6 reporting, the HF count ratio is 1.073 (1.006–1.144) per five official cold days, $q = 0.192$. All four uncertainty constructions exclude 1. The pre-2020 estimate is stronger: 1.113 (1.053–1.176). Official cold days fall almost entirely in December–February (141 of 145), so after calendar-month indicators the remaining contrast is between winters, not summer versus winter. HF hot nights are 1.003 (0.976–1.031). HF very hot days are 0.995 (0.963–1.028). A heat-only reading of the paper would miss the more coherent residual. That sentence is already in the live Discussion. Afterload is the reason it is scientifically available as interpretation, not as a new estimate.

Afterload is the pressure the heart pumps against. Cold exposure can raise sympathetic outflow, peripheral resistance, and therefore afterload. A heart that is already in failure has little reserve against that load. Ikäheimo (2018) reviews cardiovascular disease, cold exposure, and exercise in those terms: sympathetic vasoconstriction, raised peripheral resistance, increased cardiac workload. The review's own language is that a heart-failure patient has little physiological reserve, and that a failing heart may not be able to compensate. Li, Wu, Xu and colleagues (2026) synthesise cold as a sympathomimetic stimulus with afterload augmentation and decompensation risk in cardiovascular disease, including heart failure. Neither paper is this extract. Neither measures first hospitalisation after a first HF diagnosis without recorded admission cause. Neither supplies a monthly count ratio per five official cold days. They earn their place as the physiological reason the HF–cold residual is more coherent, across standard-error constructions and in the pre-2020 window, than the CHD–hot-night residual.


Goggins and Chan (2017) remain the local epidemiological neighbour, not a magnitude to import. Their daily public-hospital heart-failure admissions during 2002–2011 had a cumulative relative risk of 2.63 (2.43–2.84) comparing 11 °C with 25 °C over lags to 23 days. That is a daily, cause-coded, earlier-decade estimand. It is not 1.073 per five official cold days in a first-event series without admission cause. The live Discussion already states the comparison is between questions, not between magnitudes. This review does not move 2.63 into Table 2.

The identification of the cold-day contrast is weather-side, and it is already drawn. Remaining variation after calendar-month indicators is between-year winter. January 2021 had 11 official cold days; January 2019 had one. The model is not discovering that winter is colder than July. It is using differences among Januarys, Februarys, and Decembers. Infection, indoor heating, and care-seeking are not separated on a monthly grain. Influenza is missing for January–October 2013 and was not entered in Model 1. Those are reasons the afterload sentence stays a hedge.

The 8-df time-trend spline opens the HF cold-day interval onto 1 (1.062, 0.994–1.135). The 6-df spline does not (1.072, 1.005–1.143). That is a property of the nuisance smooth, not a myocardial duration. Physiology does not put a cliff at day 6. Afterload does not rescue the 8-df opening, and it does not convert $q = 0.192$ into multiplicity protection.


Keep the mechanism in the draft, in the live wording: cold exposure can raise afterload; a failing heart has little reserve; the present counts do not measure afterload, blood pressure, or infection. Do not name a ventricular phenotype. This cohort has no recorded HF subtype. Do not add a haemodynamic tour to the manuscript. Do not treat concordance across standard-error methods as a freeze.

Table 2: twelve Model 1 fits. Minimum $q = 0.192$.



HA_APPROVED_AGGREGATE. Live Table 2. Teal: interval excludes 1. Grey: includes 1. All twelve $q > 0.19$.

Uncertainty ladder: CHD is construction-dependent; HF is not



HA_APPROVED_AGGREGATE. Live Table 3. $q = 0.192$ for both contrasts. Do not pick an SE because it excludes 1.

Cold-related afterload is a haemodynamic hypothesis, not an estimate



This panel measures none of the middle or right boxes.

SCHEMATIC. Afterload = the pressure the heart pumps against. Not blood pressure, infection, or phenotype in this extract.

Point 3. A confinement study of seven men (drop-candidate)

This is a drop-candidate, not a keep-mechanism. The live Discussion currently names a confinement study of seven men that reported higher nocturnal core temperature after hot nights. The recommendation is to drop or collapse that adjacency, after Hogan has seen it on the bullet list. It is not a finding of this extract, and it is not a neighbour of the estimand.

Ioannou, Tsoutsoubi, Mantzios and colleagues (2024) confined seven young healthy men for ten days. Nights 1–3 and 7–10 were 22.3 °C. Nights 4–6 were 26.3 °C. Mean nocturnal core temperature was 0.2 °C higher during *and after* the three hot nights than before them, more so in the first two hours of sleep (+0.3 °C). Total sleep time fell by about 14 minutes per 0.1 °C rise in overnight core temperature. The protocol also used hot days (35.4 °C) during the heatwave block. The fact that matters for a non-recovery hypothesis is the *after* period: nights 7–10 at 22.3 °C still showed elevated nocturnal core temperature. That is laboratory support for failed overnight recovery across consecutive hot nights in young men in a chamber.

None of those design facts is this project. The participants are not people with type 2 diabetes or hypertension. They are not in Hong Kong. The outcome is not first CHD hospitalisation. The exposure is not an official monthly count of $T_{min} \geq 28$ °C at Headquarters. Core temperature is measured. Sleep is measured. The sample is seven. Consecutive nights are the experimental unit. Equation (1) has 132 months, no core temperature, no sleep, and a monthly sum that cannot distinguish consecutive nights from nights scattered through the month.

The paper is therefore the wrong object to sit beside 1.022. It can be cited in a physiology notebook as a reason consecutive-night designs exist in the laboratory. It cannot be asked to explain a territory-month count ratio.

Leaving the seven-man study in the live adjacency invites a reader to treat chamber core temperature as if it travelled to this extract. It also reads as an indoor, personal-environment claim. Hogan's 24 August comment already asked not to raise unmeasured housing and behaviour. The live body already omits the words housing and medication. A chamber study of nocturnal core temperature can still violate that comment by implication.

The drop does not delete Chevance or the overnight-recovery hypothesis. Those stay. What goes is the specific $n=7$ adjacency that the live file currently uses as a second exhibit after Chevance. If a consecutive-night physiology is ever tested in this programme, it needs a daily, cause-recorded series. The clean-room monthly-outcome/daily-exposure recovery already failed its 500-replicate calibration (null Type I error up to 0.150; minimum coverage 0.840; maximum non-null relative bias 32.8; maximum moderate false-sign rate 0.808). No real daily coefficient is reported. Ioannou does not replace that refusal.

Hogan asked to look at the paper. He did not ask to vote it onto Table 2. The sendable list already carries the DOI. This review is the longer account of why it is listed as “thinking of dropping, not adding.” Do not cut the live file until he has seen the bullets. Do not walk the seven-man protocol in the room as if it were a Hong Kong result.

Ioannou 2024 is a chamber study. This extract is a territory-month count.

Ioannou et al. 2024

n = 7 young men
10-day confinement
Nights 4–6 at 26.3 °C
+0.2 °C nocturnal core T
during and after

This extract

132 territory-months
T2D/HTN first CHD/HF
No admission cause
Outdoor official counts
No core temperature

SCHEMATIC. Protocol numbers from Ioannou et al. 2024; not a coefficient in Table 2.

Point 4. Bedroom temperature and heart-rate variability (drop-candidate)

This is the second drop-candidate. The live Discussion currently states that bedroom temperatures above 24 °C were associated with lower heart-rate variability in older adults. That sentence is true of the cited study. It is not a measurement in this series. It reads as if the monthly file had indoor bedrooms.

O’Connor, Bach, Forbes and colleagues (2025) monitored 47 community-dwelling adults aged 65 years or older in southeast Queensland across one summer (December 2024–March 2025). Wearables recorded heart rate and heart-rate variability at night. In-home sensors recorded bedroom temperature. Relative to nights below 24 °C, bedroom temperatures of 24–26 °C, 26–28 °C, and 28–32 °C had odds ratios 1.4 (1.2–1.6), 2.0 (1.8–2.3), and 2.9 (2.5–3.4) for a clinically relevant drop in lnRMSSD. Heart rate rose. Frequency-domain indices shifted toward sympathetic dominance. Across 14,179 valid nighttime hours, median bedroom temperature was 25.9 °C. The paper is observational, Australian, and about autonomic strain in older adults’ bedrooms. It is open access. It is not first CHD hospitalisation in Hong Kong. It is not $T_{min} \geq 28\text{ °C}$ at Headquarters.


Hong Kong’s official hot night is an outdoor threshold at one station. O’Connor’s 24 °C is an indoor bedroom bin. Equating them would convert an official monthly count into a bedroom temperature. The live file does not measure indoor temperature. Hogan’s weather paragraph averages daily Headquarters observations to calendar months. That paragraph is verbatim and is not being rewritten. It cannot be asked to yield Queensland bedroom temperatures, or Hong Kong bedroom temperatures.

The paper is useful as a reason indoor night temperature is a *hypothesis* for a different design. It is the wrong exhibit beside a monthly outdoor count. Left in the live adjacency, it invites the reader to treat 1.022 as if it were an indoor autonomic finding. That is the same class of overclaim as treating Ioannou’s chamber as this cohort. It also raises, by implication, an unmeasured indoor environment after Hogan asked not to raise unmeasured housing and behaviour.

Ashe et al. (2025) again bound the story. Three studies; mixed blood pressure; no demonstration of sleep as a mediator of hospitalisation. O’Connor is one observational HRV study in older adults. It does not close a pathway from official hot nights to first CHD hospitalisation. It also does not measure hospitalisation at all.

Drop or collapse the bedroom-HRV adjacency after Hogan has seen the bullets. Keep the sentence that those studies do not measure first CHD hospitalisation in Hong Kong and do not measure indoor temperature in this series. Do not add a 24 °C indoor guideline to the manuscript. Do not evaluate WHO daytime indoor guidance from this file. Do not treat outdoor official counts as a proxy that has been calibrated to bedrooms.

Outdoor official counts do not measure bedroom temperature



Point 5. Not proposing indoor temperature from outdoor counts

This is a refusal, not a paper. The sendable list groups it with two other refusals so Hogan does not look for a missing DOI. The contribution to the project is negative and specific: the analysis will not invert Headquarters monthly counts into indoor temperature, air-conditioning use, or shelter occupancy.

The weather co-investigator's live paragraph states that meteorological data was obtained from the HKO, and that monthly variables were derived by taking the average of daily data in each calendar month. Official hot nights are days with $T_{min} \geq 28$ °C. Official very hot days are $T_{max} \geq 33$ °C. Official cold days are $T_{min} \leq 12$ °C. Those are outdoor, station-based, daily indicators summed to months. They are the exposures in Model 1. They are not bedroom temperatures, ward temperatures, or personal heat indices.

O'Connor et al. (2025) and Chevance et al. (2024) are the papers that make indoor temperature scientifically tempting. Chevance includes indoor as well as outdoor temperature among the heat exposures associated with degraded sleep. O'Connor measures bedrooms. Guo et al. (2024) show that an hourly nighttime excess metric, which is still outdoor but finer than the official flag, associates with hospitalisation when the official flag does not. None of those results licenses a two-step: take the monthly official count, declare it indoor, and interpret 1.022 as a bedroom effect. There is no transfer function in this repository from Headquarters T_{min} to a Hong Kong bedroom. Housing quality, cooling, and behaviour are unmeasured. Hogan already asked that they not be raised as if they were in the file.

WHO 2026 heat-health action-plan element 6 is reducing heat exposure at household, building, and urban scales. Hong Kong already issues indoor ventilation and air-conditioning advice and operates temporary heat shelters. The monthly panel cannot inform whether those instruments reduce indoor night temperature in the diabetes and hypertension group. It cannot count shelter nights. It cannot report air-conditioning prevalence. Mapping element 6 onto existing advice is already in the companion HHAP note. This refusal is the matching methods sentence: indoor temperature is not recovered from outdoor counts.

The same refusal applies to cold. Indoor heating in January is unmeasured. Afterload remains a hypothesis partly because indoor conditions are unknown. That is a reason to keep the hedge, not a reason to invent an indoor-cold exposure. Official cold days are $T_{min} \leq 12$ °C outdoors. They are not a measured indoor heating deficit.

A modelled indoor series built from outdoor percentiles would still be outdoor weather under another name. It would not become O'Connor's in-home sensor. It would not become a Hong Kong bedroom. The live limitations already state that all exposures were measured at a single Observatory station and applied territory-wide. That sentence is the indoor-environment bound. It does not need a new covariate.

If a later design measures indoor night temperature, or personal exposure, that will be a different paper. It will need a different DUA, a different estimand, and a different Hogan weather lock. Until then, the honest contribution is the refusal. Do not add a modelled indoor series. Do not treat the 24 °C bedroom bin as a Hong Kong threshold. Do not rewrite the HKO paragraph to imply multi-station indoor averaging. That paragraph stays as it is.

Point 6. Five extra official days is a reporting scale, not consecutive duration

This is the identification point that looks like physiology and is not. Model 1 divides the monthly official count by five, $I(\text{count}/5)$, so that e^{β} is the count ratio per five additional such days in the month. Five is a convenient scale, like reporting a 1 °C contrast rather than a 0.1 °C contrast. It is not a new weather threshold. It is not a consecutive-duration rule. The live Methods already say this. The live Discussion already says a monthly official-day total does not distinguish consecutive days from days scattered through the month.

The exhibit that makes the distinction empirical is 2018. In that year Hong Kong recorded 26 official hot nights and zero days inside a run of five or more consecutive official hot nights. A year can be heavy in official hot nights without containing a five-night spell. Across 2013–2023 the same annual table records 449 official hot nights, of which 195 (43.4%) sat inside ≥ 5 -night spells. The other 56.6% did not. Very hot days show a similar split (421 days; 183, 43.5%, inside ≥ 5 -day spells). 2018 is the opposite pattern for very hot days: 36 very hot days and 26 days inside ≥ 5 -day spells. The official-count exposure and the spell-day exposure are not interchangeable, and they do not move together year by year.


Wang and colleagues' ≥ 5 consecutive very-hot-days or hot-nights convention, and Ho and colleagues' consecutive-night mortality contrast, are local *mortality* designs. They are a reason to test consecutive nights if a daily cause-recorded series is obtained. They are not a reason to read $I(\text{count}/5)$ as that test. The Dutch meteorological heatwave (≥ 5 summer days) is a third literature. None of those objects is Table 2.

Spline degrees of freedom are a fourth object that has been fused with "5, 6, 8." Model 1 is a 4-df time-trend spline. Sensitivities are 3, 6, and 8 df, or year indicators. Those numbers describe how flexible the long-term time smooth is. They do not describe a 3-day, 6-day, or 8-day heatwave. Among specifications more flexible than Model 1, only the 6-df spline keeps both headline intervals away from 1. The 8-df spline opens HF cold days onto 1. Physiology supports nights and non-recovery. It does not put a cliff at day 6. The 6-df spline is a named robustness check. It does not replace Table 2.

The failed daily-recovery calibration is the corresponding methods limit. Monthly sums do not automatically yield daily trigger estimates in this implementation. That refusal is specific to this series, this method, and this calibration standard. It is not evidence that recovery of daily effects from aggregated outcomes fails in general. It is evidence that this paper will not report a consecutive-night coefficient from 132 monthly counts.

Not proposing consecutive duration therefore contributes in three places: the Methods sentence on the reporting scale; the Discussion sentence that a monthly total cannot identify consecutive days; and the 2018 exhibit, which is weather, already in Hogan's domain, and needs no new health model. Do not adopt a five-night warning from this file. Do not describe Figure 3 as a 5/6/8-day heatwave forest. Do not treat 2018's 26 nights as proof of a spell effect. That year had zero five-night spell-days.

Official hot-night counts are not consecutive-night spells




REAL public HKO Headquarters annual table. A spell-day is a day inside a run of ≥ 5 consecutive official hot nights.

Official-count Model 1 estimates: time-trend and window checks

● 95% CI excludes 1
● 95% CI includes 1

Rows labelled 3, 6, or 8 df change the time-trend spline, not heatwave length.
 The scale is five additional official days in the month, not consecutive days.



Health provenance: HA_APPROVED_AGGREGATE. Newey–West lag-6 95% intervals; days-in-month offset. Shaded row is Model 1 (4-df time-trend spline). Continuous-temperature fits for the same nine specifications are Supplementary Figure S6. Source: [outputs/tables/cvd_trend_depletion_sensitivity.csv](#). No new health model.

Point 7. Not evaluating Hong Kong's heat or cold warnings

This is the second policy refusal. Hong Kong issues official heat and cold warnings. The estimates do not evaluate those warnings. They do not set a five-day trigger. They do not count admissions averted. They do not say whether existing instruments work. The live Discussion already states this. The contribution of a longer note is to show *why* the monthly file cannot do the evaluation, and what it can do instead: a map.

The World Health Organization Regional Office for Europe's 2026 heat-health action-plan guidance lists eight core elements: governance; a heat-health warning system; populations at increased risk; communication; health-system resilience; reducing heat exposure; heat-health surveillance; and monitoring, evaluation and learning. Hong Kong is not a European member state. It already operates products that sit near several of those elements: the Very Hot Weather Warning; a Hot Weather Special Advisory; a Prolonged Heat Special Alert that names hot nights beside very hot days; an Extremely Hot Weather product around 35 °C; Centre for Health Protection advice that names people with heart disease or high blood pressure; temporary heat and cold shelters; and a separate Cold Weather Warning. Those are instruments. They are not, on the 2026 definition, the eight documented HHAP output files. The monthly panel cannot supply the missing outputs.

What the panel can inform is limited and already used. The CHD residual sits with official hot-night *counts*, not very-hot-day counts (0.999). The Prolonged Heat product already names nights. Department of Health materials already name heart disease and hypertension; this cohort is inside that named group, not a new named group. A heat-only HHAP reading would omit the more coherent residual, which is HF with official cold days. WHO heat guidance does not replace the Cold Weather Warning.


What the panel cannot inform is the evaluation itself. There is no warning-day treatment indicator in Model 1. There is no daily, cause-recorded outcome. Admission cause is absent, so an event is a first hospitalisation after a first diagnosis, not a cardiac-caused admission on a warning day. Monthly first-event counts are not operational surveillance for activation (element 7). They are not a monitoring-evaluation-learning framework (element 8). Law and colleagues' daily study of mortality and emergency calls during extremely hot weather is a neighbouring estimand and is not this file.

The CHD serial-correlation problem (lag-1 Pearson autocorrelation 0.508 in the hot-night model) and the open multiplicity problem (minimum $q = 0.192$) are further reasons not to promote 1.022 into a statement that a warning should be tightened or relaxed. The HF residual is more coherent and still q -unprotected. Evaluation, if it is ever done, needs a predeclared contrast, warning-day exposure, and cause-recorded daily outcomes. Those data are not this extract.

Not proposing an evaluation therefore protects the paper's actual contribution: identification of what monthly first-event counts can and cannot settle, with both heat and cold in one design. Keep the mapping paragraph in the live file. Keep the sentence that the estimates do not evaluate warnings. Do not add a score for VHWW. Do not report admissions averted. Do not send Hogan a second HHAP document; the bullets already say the refusal in one clause.

Cold days are concentrated in winter

141 of 145 days occurred in December–February (Dec 40, Jan 54, Feb 47); the remaining 4 occurred in March



Environmental provenance: REAL public Hong Kong Observatory Headquarters data; cold day = daily minimum temperature ≤12°C. Identification: remaining variation after calendar-month indicators is between-year winter, not summer-versus-winter.

What this packet does not do

It does not replace the sendable bullets. It does not rewrite Hogan's HKO paragraph. It does not cut the live Discussion. It does not freeze a confirmatory claim. It does not fit Models 2 or 3. It does not invent a stroke series. Physiology remains interpretation. The next human step, if Monday's meeting happens, is still the short list with DOIs, not this PDF.